

# Cost-based data integration with disjunction

Lucas WILLEMS

5 September 2017

Informally, the PDQ system (Proof-Driven Querying) works as follow:

- inputs:
  - a *schema* consisting of *relations*.
  - the data described by the schema is not freely available to a user. Instead it can be obtained only via a set of *access methods*. Each access method can expose some data for a given cost.
  - a user *query* written over the schema.
- output: if possible, the **lowest-cost** plan that answers the query; otherwise, nothing.

**My intership purpose** was to add support for *disjunction* in the constraints of these schemas in order to support more realistic relationships. This would add *union* at the plan level.

A *database* is given by:

- a *schema*  $Sch$  (the “organisation”):
  - a finite collection of *relations*. A *relation* is a *relation name* associated with an *arity*, i.e. a positive number.
  - a finite collection of *schema constants* denoted  $c\vec{st}$ .
  - a collection of *integrity constraints*, i.e. sentences in some logic.
- an *instance* (the “data”) for this schema  $Sch$  that is a collection of *facts* such that all integrity constraints of  $Sch$  are satisfied. A *fact* is the association of a relation  $R$  with a tuple  $\vec{c}$  of the proper arity.

To request for information from a database, PDQ uses *conjunctive queries* of the form:  $Q(\vec{x}) = \exists \vec{y} \bigwedge_{i=1}^n A_i(\vec{x}, \vec{y}, c\vec{st})$  where  $A_i$  are relational atoms.

- a *binding* of  $Q$  is a mapping from its free variables to some values.
- an *homomorphism* of  $Q$  in an instance  $I$  is a binding that witnesses that  $Q$  holds in  $I$ .
- the *output* of  $Q$  over  $I$ , denoted  $\llbracket Q \rrbracket_I$ , is the collection of the homomorphisms of  $Q$  over  $I$ .

These constraints are built up from *atoms*, which can be either:

- *relational atoms* of the form  $R(\vec{t})$ , where  $R$  is a relation name and each  $t_i$  in  $\vec{t}$  is a constant or a variable;
- *equality atoms* of the form  $t_i = t_j$  with  $t_i, t_j$  a constant or a variable.

Below,  $\varphi$  will denote to a conjunction of atoms and  $\rho$  and  $\rho_i$  to conjunctions of relational atoms.

Before starting my internship, only support for:

- *equality-generating dependencies* (EGDs):  $\forall \vec{x} [\varphi(\vec{x}) \rightarrow x_i = x_j]$  where  $x_i$  and  $x_j$  are distinct variables.
- *tuple-generating dependencies* (TGDs):  $\forall \vec{x} [\varphi(\vec{x}) \rightarrow \exists \vec{y} \rho(\vec{x}, \vec{y})]$ .

After my internhsip, additional support for *disjunctive TGDs* (DTGDs):

$$\forall \vec{x} \left[ \varphi(\vec{x}) \rightarrow \exists \vec{y} \bigvee_{i=1}^n \rho_i(\vec{x}, \vec{y}) \right]$$

### Example

A possible schema for a library:

- relations: Books(id, title), Members1(id, name), Members2(id, name), Loans(bid, mid).
- schema constants: Dumbo.
- integrity constraints:
  - some EGDs to make the IDs unique.
  - $\forall x y [\text{Loans}(x, y) \rightarrow \exists z \text{Books}(x, z)]$ ;
  - $\forall x y [\text{Loans}(x, y) \rightarrow \exists z \text{Members1}(y, z) \vee \text{Members2}(y, z)]$ .

With some instance:

{Books(1, Dumbo), Books(2, Aladdin), Members(1, Michael), Members(2, Efi), Loans(1, 1), Loans(2, 1)}.

We may request for:

$$Q(x) = \exists y z \text{Books}(y, \text{Dumbo}) \wedge \text{Loans}(y, z) \wedge \text{Members}(z, x).$$

# Database knowledge

## Query containment

A common database problem is called *query containment with constraints*. Given two queries  $Q$  and  $Q^*$  and a conjunction of integrity constraints  $\Sigma$ , we want to determine if  $Q \wedge \Sigma \models Q^*$ , i.e. if, for every instance  $I$  that satisfies  $\Sigma$ ,  $\llbracket Q \rrbracket_I \subseteq \llbracket Q^* \rrbracket_I$ .

For  $Q$  and  $Q^*$  conjunctive queries and  $\Sigma$  a conjunction of TGDs, *the chase* can be used to prove these entailments:

- ① start with the *canonical database* of  $Q$  ( $\text{CanonDB}(Q)$ ): the instance with a fact  $R(c_1, \dots, c_n)$  for each atom  $R(x_1, \dots, x_n)$  of  $Q$ .
- ② iteratively derive consequences with the constraints  $\Sigma$ , i.e repeat:
  - ① select a *trigger* for a TGD  $\delta = \forall \vec{x} [\varphi(\vec{x}) \rightarrow \exists \vec{y} \rho(\vec{x}, \vec{y})]$ , i.e. a tuple  $\vec{e}$  such that  $\varphi(\vec{e})$  holds.  $e$  is *active* if there is no  $\vec{f}$  such that  $\rho(\vec{e}, \vec{f})$  holds.
  - ② *fire the rule*, i.e. add the facts that make  $\rho(\vec{e}, \vec{f})$  true.
- ③ stop and declare success if we have reached an instance that has a “match” for  $Q^*$ , i.e. if  $x_i \mapsto c_i$  is an homomorphism of  $Q^*$ .

A *chase sequence following a set of TGDs  $\Sigma$*  is a sequence of instances, called *chase configurations*, where a configuration is related to its successor by a rule firing of a constraint in  $\Sigma$ .

In PDQ, several chases are implemented but they only execute the step 2.

# The global working of PDQ

## Access methods and plans

PDQ takes as input:

- an *access schema* that consists of:
  - a schema with, originally, TGDs as integrity constraints.
  - for each relation  $R$ , a collection of *access methods*. Each access method is associated with an *access cost* and a collection of *input positions* of  $R$ , i.e. a collection of numbers between 1 and  $\text{arity}(R)$ .
- a conjunctive query  $Q$  written over the input schema.

and then outputs, if possible, a plan that answers  $Q$ .

- A *plan* is a sequence of access and middleware query commands.
- It *answers* the query if for every instance  $I$  satisfying the constraints of the input access schema, the output of the plan on  $I$  is the same as the output of  $Q$ .



# The global working of PDQ

## Query containment reformulation (1/2)

PDQ's plan search algorithms start by transforming the input access schema  $Sch$  into a *forward accessible schema for  $Sch$*   $AcSch(Sch)$ :

- the constants are those of  $Sch$ .
- the relations are:
  - *original relations*, i.e. the relations of  $Sch$ .
  - *InfAccCopy relations*, i.e. a copy of each relation  $R$  called  $InfAccR$ .
  - a unary relation  $accessible(x)$ .
- the constraints are:
  - *original integrity constraints*, i.e. the integrity constraints of  $Sch$ .
  - *InfAccCopy integrity constraints*, i.e. a copy of each of the integrity constraints of  $Sch$ , with each relation  $R$  replaced by  $InfAccR$ .
  - *accessibility axioms*, i.e. for each access method on relation  $R$  of arity  $n$  with input positions  $j_1, \dots, j_m$  (universal quantifiers omitted):

$$\left[ R(x_1, \dots, x_n) \wedge \bigwedge_{k=1}^m accessible(x_{j_k}) \rightarrow InfAccR(x_1, \dots, x_n) \wedge \bigwedge_{j=1}^n accessible(x_j) \right].$$

# The global working of PDQ

## Query containment reformulation (2/2)

At this point, finding plans comes down to exploring the space of proofs of this entailment:

$$Q \wedge \Gamma \models \text{InfAcc}Q$$

where  $\Gamma$  is the conjunction of the constraints of  $\text{AcSch}(\text{Sch})$  and  $\text{InfAcc}Q$  is a query obtained from copying  $Q$  and replacing each relation  $R$  by  $\text{InfAcc}R$ .

In PDQ, several plan search algorithms are implemented. **They use the chase** to explore the space of proofs, represented by a tree of chase configurations.

Depending on the programs, the space of proofs is more or less explored.

# My work on PDQ

## My purpose

My purpose: allowing input access schemas to contain DTGDs.

What differs from using only TGDs?

- Firing a TGD  $\forall \vec{x} [\varphi(\vec{x}) \rightarrow \exists \vec{y} \bigwedge_{i=1}^n A_i(\vec{x}, \vec{y})]$  for a trigger  $\vec{e}$  on an instance  $I$  adds  $\{A_1(\vec{e}, \vec{f}), \dots, A_n(\vec{e}, \vec{f})\}$  to  $I$ , with  $\vec{f}$  a tuple of fresh values.

It is **deterministic**, we use **chase configurations**.

- Firing a DTGD  $\forall \vec{x} [\varphi(\vec{x}) \rightarrow \exists \vec{y} \bigvee_{i=1}^n \bigwedge_{j=1}^n A_{i,j}(\vec{x}, \vec{y})]$  for a trigger  $\vec{e}$  on an instance  $I$  adds either  $\{A_{1,1}(\vec{e}, \vec{f}), \dots, A_{1,n}(\vec{e}, \vec{f})\}$  to  $I$  or ... or  $\{A_{n,1}(\vec{e}, \vec{f}), \dots, A_{n,n}(\vec{e}, \vec{f})\}$  to  $I$ .

It is **not deterministic**, I introduced the notion of **disjunctive chase configurations**, i.e. collections of chase configurations.

This implies to modify:

- all PDQ's chases;
- all PDQ's plan search algorithms.

# My work on PDQ

## The restricted chase with DTGDs - Algorithm

One of PDQ's chases is the *restricted chase*, *restricted* meaning that only *active* triggers are selected.

Here is the DTGD version of this algorithm:

```
input: A disjunctive chase configuration disConfig ;  
        A collection of DTGDs  $\Sigma$  ;  
while there is an active trigger for a DTGD  $\lambda$  on a configuration  
      config of disConfig do  
    Choose such a trigger ;  
    Remove config from disConfig ;  
    foreach TGD  $\delta$  of the splitted DTGD  $\lambda$  do  
      newConfig := a copy of config ;  
      Fire the rule of  $\delta$  on newConfig ;  
      Add newConfig to disConfig ;  
    end  
  end
```

# My work on PDQ

## The restricted chase with DTGDs - Example

### Example

Parameters (universal quantifiers omitted):

- $\text{disConfig} = \{C_0\}$  where  $C_0 = \{R(c)\}$ ;
- $\Sigma = \{\lambda_1, \lambda_2\}$  with  $\lambda_1 = R(x) \rightarrow S(x) \vee T(x)$ ,  $\lambda_2 = S(x) \rightarrow T(x) \vee U(x)$ .

Execution:

- 1 Active trigger (c) for  $\lambda_1$  on  $C_0$ .  
 $C_0$  replaced by  $C_1 = \{R(c), S(c)\}$  and  $C_2 = \{R(c), T(c)\}$ .  
Hence,  $\text{disConfig} = \{C_1, C_2\}$ .
- 2 Active trigger (c) for  $\lambda_2$  on  $C_1$ .  
 $C_1$  replaced by  $C_3 = \{R(c), S(c), T(c)\}$  and  $C_4 = \{R(c), S(c), U(c)\}$ .  
Hence,  $\text{disConfig} = \{C_2, C_4, C_5\}$ .
- 3 No other active trigger for a configuration of  $\text{disConfig}$ .  
Finally,  $\text{disConfig} = \{\{R(c), T(c)\}, \{R(c), S(c), T(c)\}, \{R(c), S(c), U(c)\}\}$ .



# My work on PDQ

## Problems encountered and all the chases with DTGDs

I encountered several problems:

- storing several instances at the same time in a database: use of an `InstanceID`.
- fixing bugs in the code after creating some unit tests:
  - hard;
  - time-consuming;
  - require the use of debugging tools.

Then, I added support for DTGDs in all the chases:

- an abstract class called `Chaser` that contains the common part
- a class for each different chase that:
  - inherits this abstract class;
  - defines a `reason` method, containing the chase strategy.

# My work on PDQ

## The naive plan search program with DTGDs - Algorithm

One of PDQ's plan search algorithms is the *naive* one. It explores the full space of proofs of  $Q \wedge \Gamma \models \text{InfAcc}Q$ .

For the TGD version, the space of proof corresponds to a tree of chase configurations. In the DTGD version, it is represented by a **list of trees of disjunctive chase configurations**:

- ① chase  $\text{disConfig} = \{\text{CanonDB}(Q)\}$  with the **original integrity constraints**. For each chase configuration  $C$  obtained, create a tree with the disjunctive chase configuration  $\{C\}$  as root.
- ② iteratively expose facts in the different trees, i.e. repeat:
  - ① select an active trigger for an **accessibility axiom**  $\delta$  on a disjunctive configuration  $\text{disC}$ .
  - ② chase  $\text{disC}$  with  $\delta$  and then with  $\text{InfAccCopy}$  **integrity constraints**.
  - ③ add this new configuration  $\text{newDisC}$  to  $\text{disC}$ 's children, check if  $\text{InfAcc}Q$  has a match on  $\text{newDisC}$ , update the best plan of the tree.
- ③ stop when there is no such trigger anymore, return the union of the best plan found in each tree.


# My work on PDQ

## The naive plan search program with DTGDs - Example (1/2)

### Example

#### Parameters:

- an access schema composed of:
  - 4 relations:
    - $R$  and  $T$  of arity 1 with no access method;
    - $A1$  of arity 1 with a free access of cost 10;
    - $A2$  of arity 1 with a free access of cost 20.
  - 4 integrity constraints  $\Gamma$  (universal quantifiers omitted):  
 $\lambda_1 = (R(x) \rightarrow A1(x) \vee A2(x))$  /  $\lambda_2 = (A1(x) \rightarrow R(x) \vee T(x))$  /  
 $\lambda_3 = (A2(x) \rightarrow R(x) \vee T(x))$  /  $\lambda_4 = (T(x) \rightarrow R(x))$ .
- a query  $Q(x) = R(x)$ .





# My work on PDQ

## The naive plan search program with DTGDs - Example (2/2)

### Example

#### Execution:

- ① Chase  $\{\text{CanonDB}(Q)\}$  i.e.  $\{\{R(c)\}\}$  with the original integrity constraints what gives  $\{C_1, C_2\}$  with  $C_1 = \{R(c), A1(c)\}$ ,  $C_2 = \{R(c), A2(c)\}$ . Create trees  $t_1$  with  $\{C_1\}$  as root and  $t_2$  with  $\{C_2\}$  as root.
- ② Search the best plan for  $t_1$ :
  - ① Active trigger (c) for the accessibility axiom  $\delta = A1(x) \rightarrow \text{InfAcc}A1(x)$  on  $C_1$ .
  - ② Chase  $\{C_1\}$  with  $\delta$  and then with  $\text{InfAccCopy}$  integrity constraints:
$$\text{newDisC} = \{\{R(c), A1(c), \text{InfAcc}A1(c), \text{InfAcc}R(c)\}, \\ \{R(c), A1(c), \text{InfAcc}A1(c), \text{InfAcc}T(c), \text{InfAcc}R(c)\}\}$$
  - ③ Add  $\text{newDisC}$  to  $t_1$ ,  $\text{InfAcc}Q(x) = \text{InfAcc}R(x)$  has a match in  $\text{newDisC}$ ,  $t_1$  has a new best plan that  $\text{newDisC}$ 's best plan of cost 10.
- ③ This is the same for  $t_2$ , but with a best plan of cost 20.
- ④ Return the global best plan that has a cost of 30.

# My work on PDQ

Problems encountered and all the plan search programs with DTGDs

I encountered several problems:

- before:
  - defining the inferred accessible DTGDs: double inheritance not possible in Java.
  - modifying the “parser” of the `schema.xml` file.
- during: misunderstanding of how the algorithm worked.
- after: debugging the code.

Then, I added support for DTGDs in all the chases:

- mostly, adapting the definitions of *dominance* and *equivalence* to disjunctive chase configurations.

# My work on PDQ

## DLV as an alternative for saturation and success checking

Michael Benedikt suggested to add support for DLV, a tool that natively handles disjunctive logic programming, as an alternative for:

- saturation, i.e. adding the InfAccCopy consequences.  
`DLV -silent {facts} {rules}`
- success checking, i.e. checking if InfAccQ has a match.  
`DLV -silent -cautious {facts} {rules} {query}`

Therefore, I had to:

- create a **DLV chase configuration**, i.e. an abstraction of a collection of DLV models.
- use the DLV reader and printer made by Efthymia Tsamoura to write and parse DLV facts and rules.
- create a DLV mapping parser.
- create an abstraction of DLV, i.e. some classes to abstract the DLV binary and its input files.

This internship helped me to:

- discover the field of databases by:
  - reading from Dr. Benedikt's notes and the book "Generating Plans from Proofs";
  - working on PDQ;
  - attending to conferences organized by the department of computer science of the University of Oxford.
- continue to develop my coding skills:
  - a new language (Java);
  - a new IDE (Eclipse) and its debugging tool.
- learn to work on a big project and deal with code written by others.
- begin to enter the world of research by talking everyday with Dr. Benedikt, the postdocs and the PhD students.